
Massive concurrent modifications in web app

Volodymyr Mykhailiyk
Anton Mishchuk



Lviv, October 2013

Outstanding RoR projects

- [Basecamp](#)
- [GitHub](#)
- [Groupon](#)
- etc, etc

one sandbox
per user



Massive multi-player online games



15000 requests per minute
“Bosses” have 30.000
modifications per hour

one sandbox for all users
massive concurrent modifications



Simple Example

MMO Cookie-clicker? Meet Cookie-hunter !!!

Cookie HunterSign Out

Cookies

Get 240 cookies per click

Regenerate 855 cookies per 10 sec

Steal 1 cookies per click

Stockpile saves 20000 cookies forever

Cookies: 126875

StealBucket: 26354

Users

Text

- test1@test.com (5337654)
- mega.hunter@cookies.com (194969)

Bonuses

Active Bonuses

Stove: * 51

Plus Click: * 16

Double Click: * 6

Grand Mother: * 69

Bread Plate: * 4

Available bonuses

Plus Click (1): Price: 17000

Double Click (*2): Price: 245000

Cookie Box (1000): Price: 1000

Bread Plate (5000): Price: 25000

Grand Mother (5): Price: 21000

Stove (10): Price: 26000

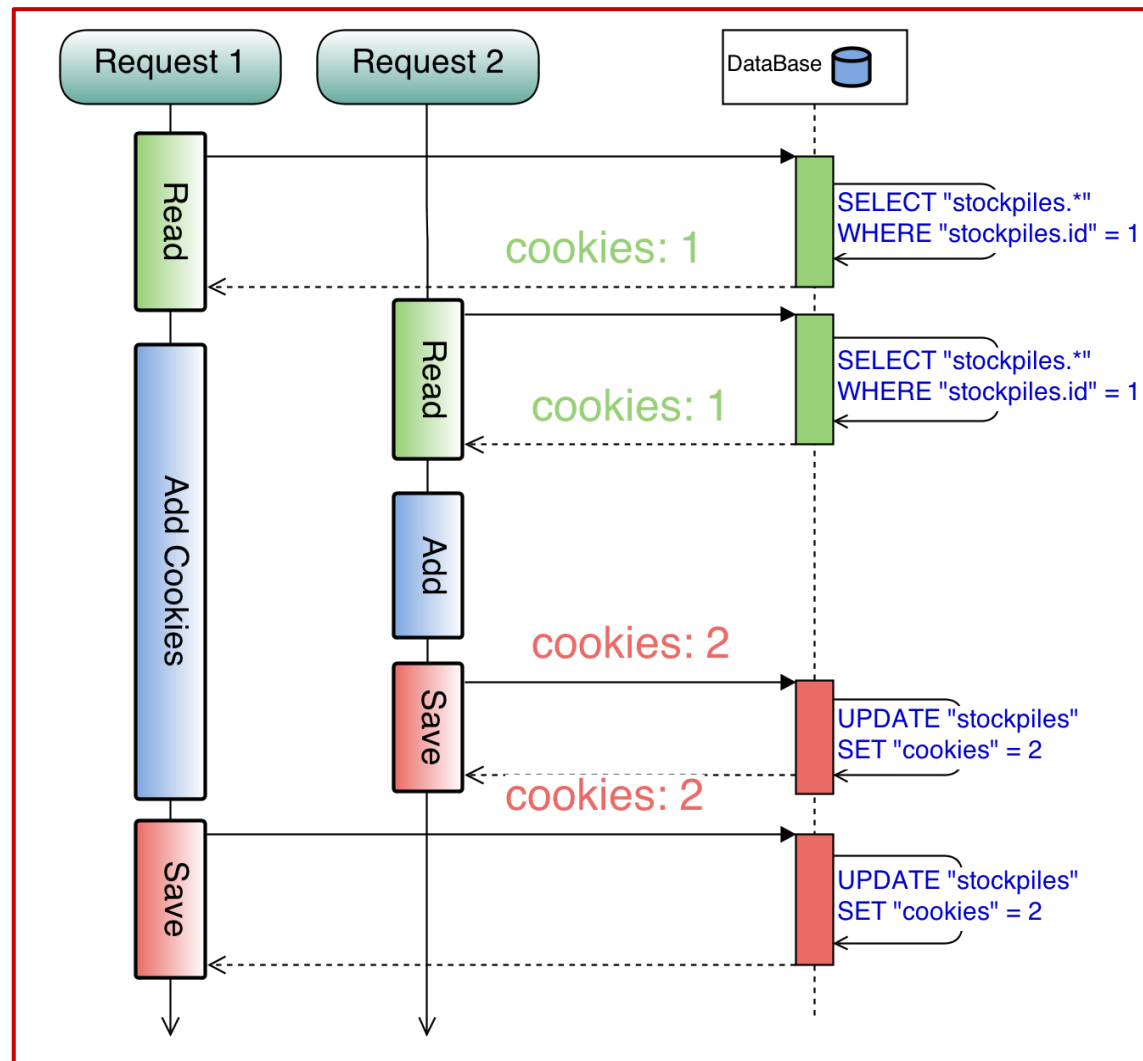
Trick (1): Price: 1000

Cheat (5): Price: 3000

Source code on GitHub

Problem 1

Race conditions



Rails counters

Model.update_counters(id, counters)

[activerecord/lib/active_record/counter_cache.rb](#)

```
UPDATE "stockpiles" SET "cookies" = COALESCE("cookies", 0) + 5
```

cons:

- uncomfortable interface
 - saves record after being called
-

Delta-attributes

The idea - instead of

```
UPDATE "stockpiles" SET "cookies" = 10
```

make this

```
UPDATE "stockpiles" SET "cookies" = "cookies" + 1
```

gem [ar-deltas](#), gem [delta_attributes](#)

```
1 class Stockpile < ActiveRecord::Base
2
3   delta_attributes :cookies
4
5 end
```

ActiveRecord::Persistence monkey-patching

Set delta (new value minus original value)
instead of value

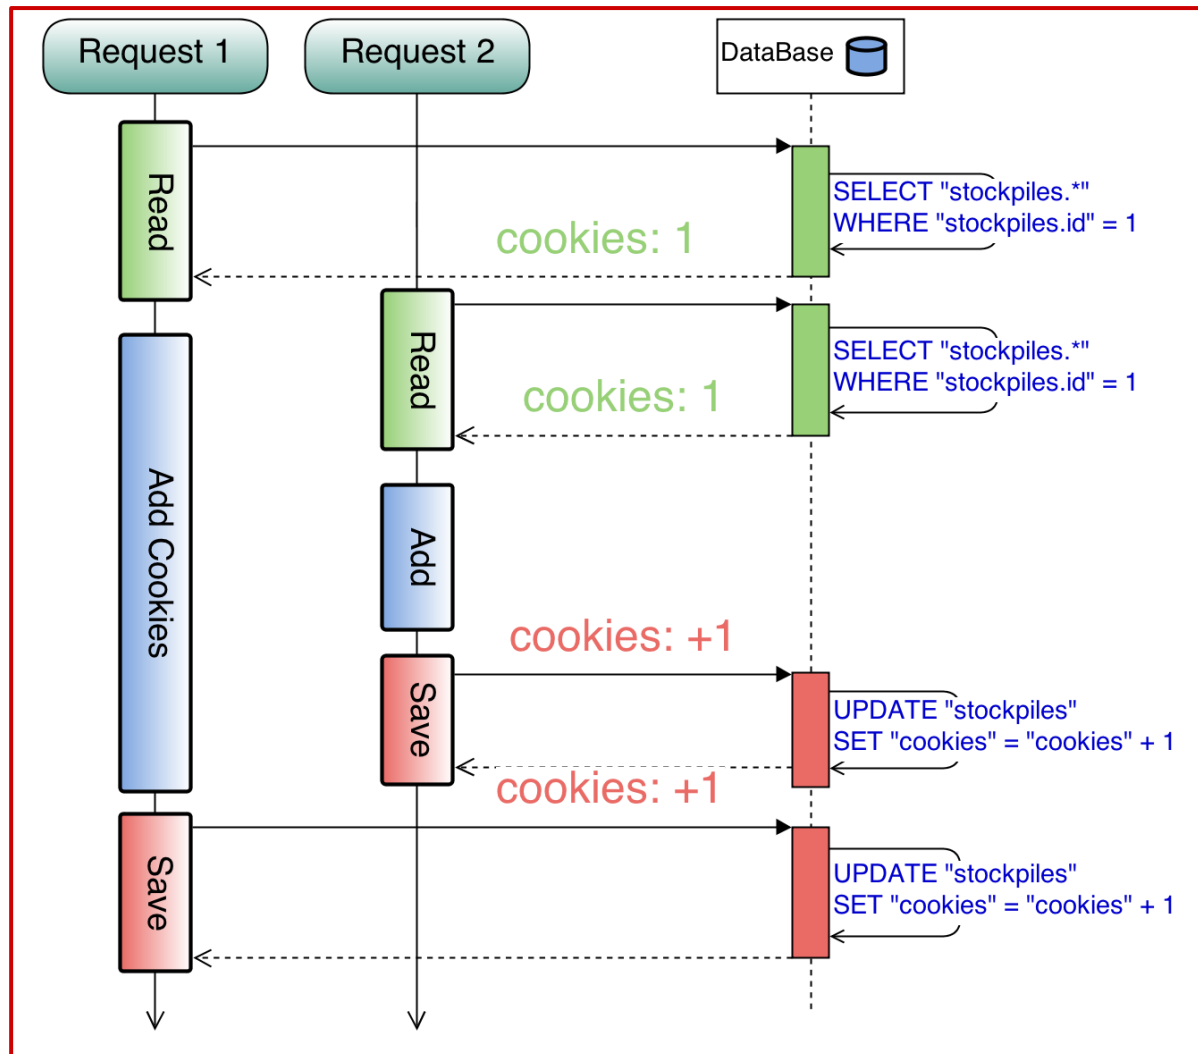
```
1  module ActiveRecord::Persistence
2
3      def update_record(attribute_names = @attributes.keys)
4          #.....
5          # begin of monkey patching code
6              v = value
7              if self.class.respond_to?(:delta_attributes)
8                  && self.class.delta_attributes.include?(attr.name)
9                  && @changed_attributes.include?(attr.name)
10
11                  v = value - @changed_attributes[attr.name]
12              end
13              [real_column, v]
14          # end
15          #.....
16      end
17
18  end
```


Arel::Visitors::ToSql monkey-patching

Modifying SQL statement

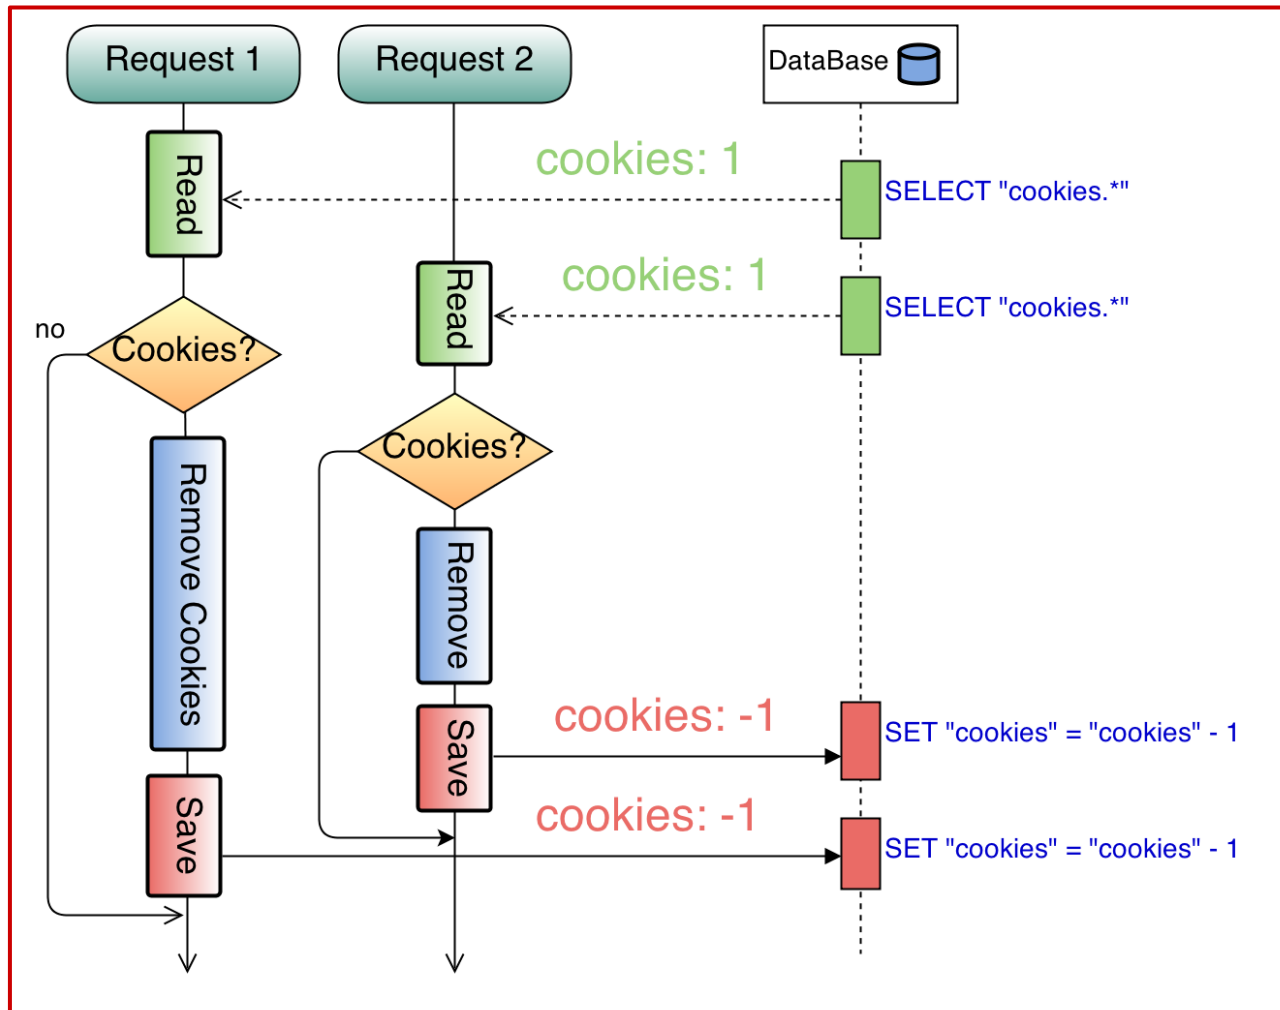
```
1  class Arel::Visitors::ToSql
2    def visit_Arel_Nodes_Assignment(o)
3      # begin monkey-patching
4      if o.left && o.left.expr && o.left.expr.relation \
5        && o.left.expr.relation.engine \
6        && o.left.expr.relation.engine.respond_to?(:delta_attributes) \
7        && o.left.expr.relation.engine.delta_attributes.include?(o.left.name)
8        l = visit o.left
9        "#{l} = #{l} + #{visit o.right}"
10     else #end
11       right = quote(o.right, column_for(o.left))
12       "#{visit o.left} = #{right}"
13     end
14   end
15 end
```

Delta attributes in action



Problem 2

Concurrency with conditions

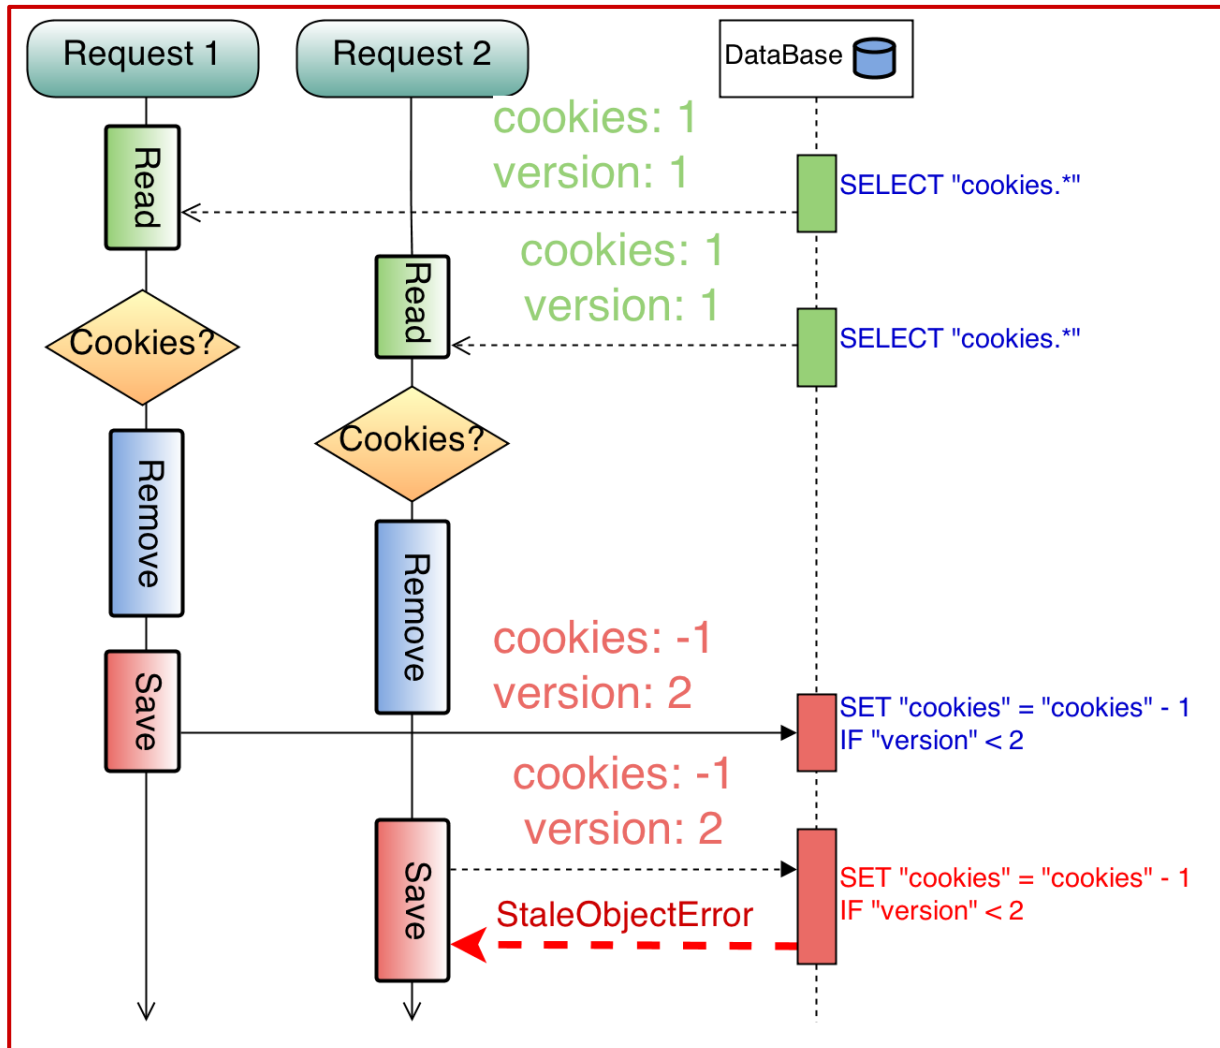


Problem 2

Concurrency with conditions

I went to the dark side
and they DID NOT
have cookies

Optimistic locking



Optimistic locking

Pros:

- no deadlocks
- no overwritten data

Cons:

- not scalable
- wasted processor time (for locked results)

Links:

- [Rails Optimistic Locking](#)
 - [Optimistic concurrency control](#)
-

Pessimistic locking

Database handles the lock

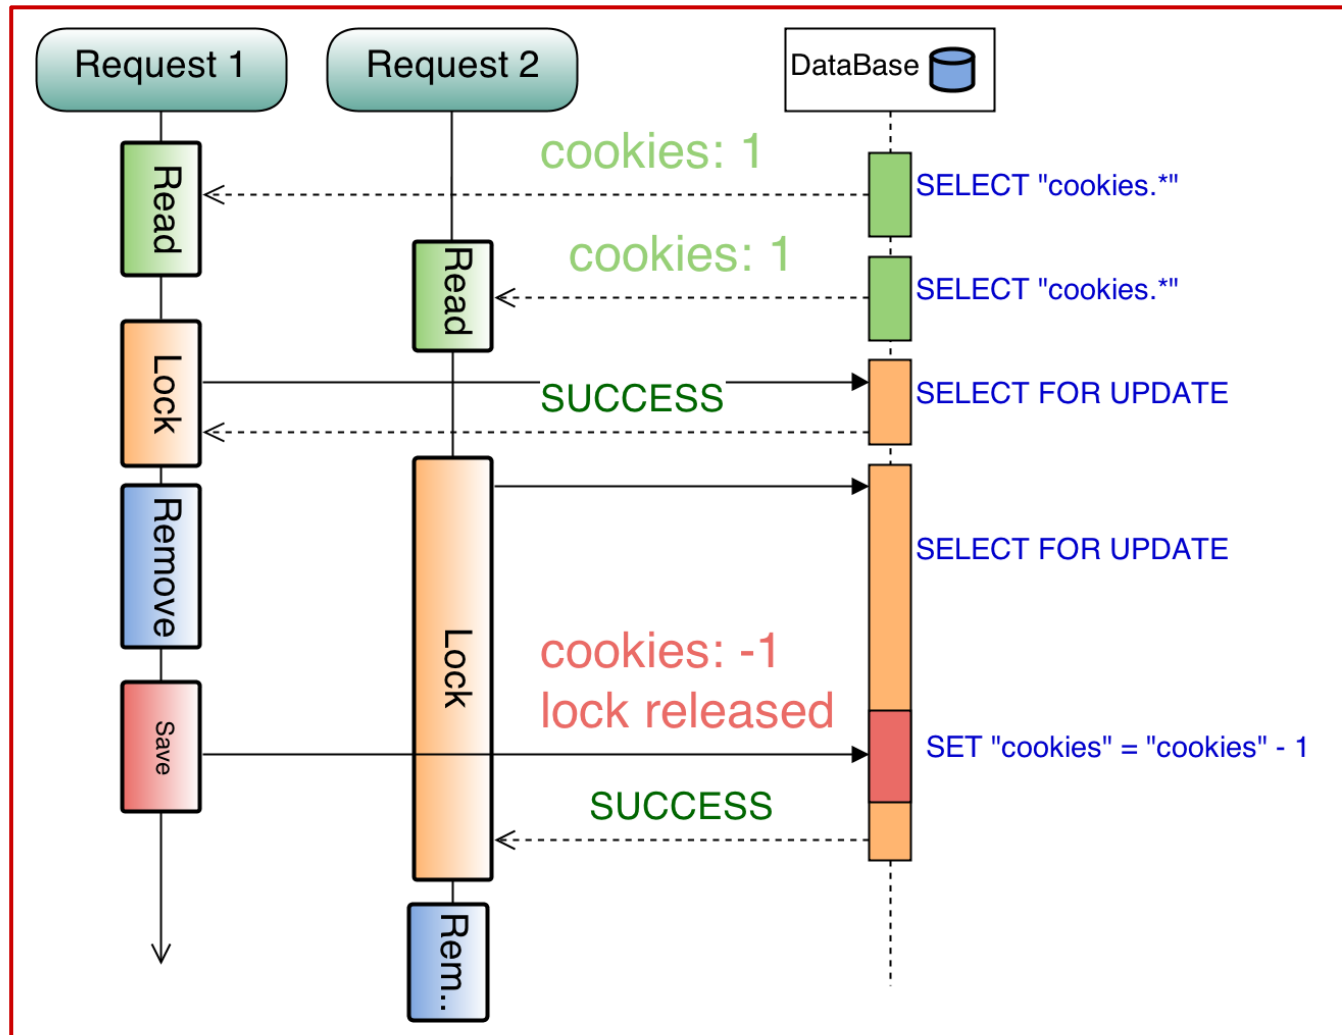
SQL Lock modes:

- **FOR_UPDATE** (mutex mode)
- **SHARE_MODE** (allow read)
- **FOR_UPDATE NOWAIT** (raise exception, even for read)

Links:

- [Rails Pessimistic Locking](#)
 - [MySql Locking](#)
 - [Postgres Locking](#)
-

Pessimistic locking (mutex)



Pessimistic locking (mutex)

Pros:

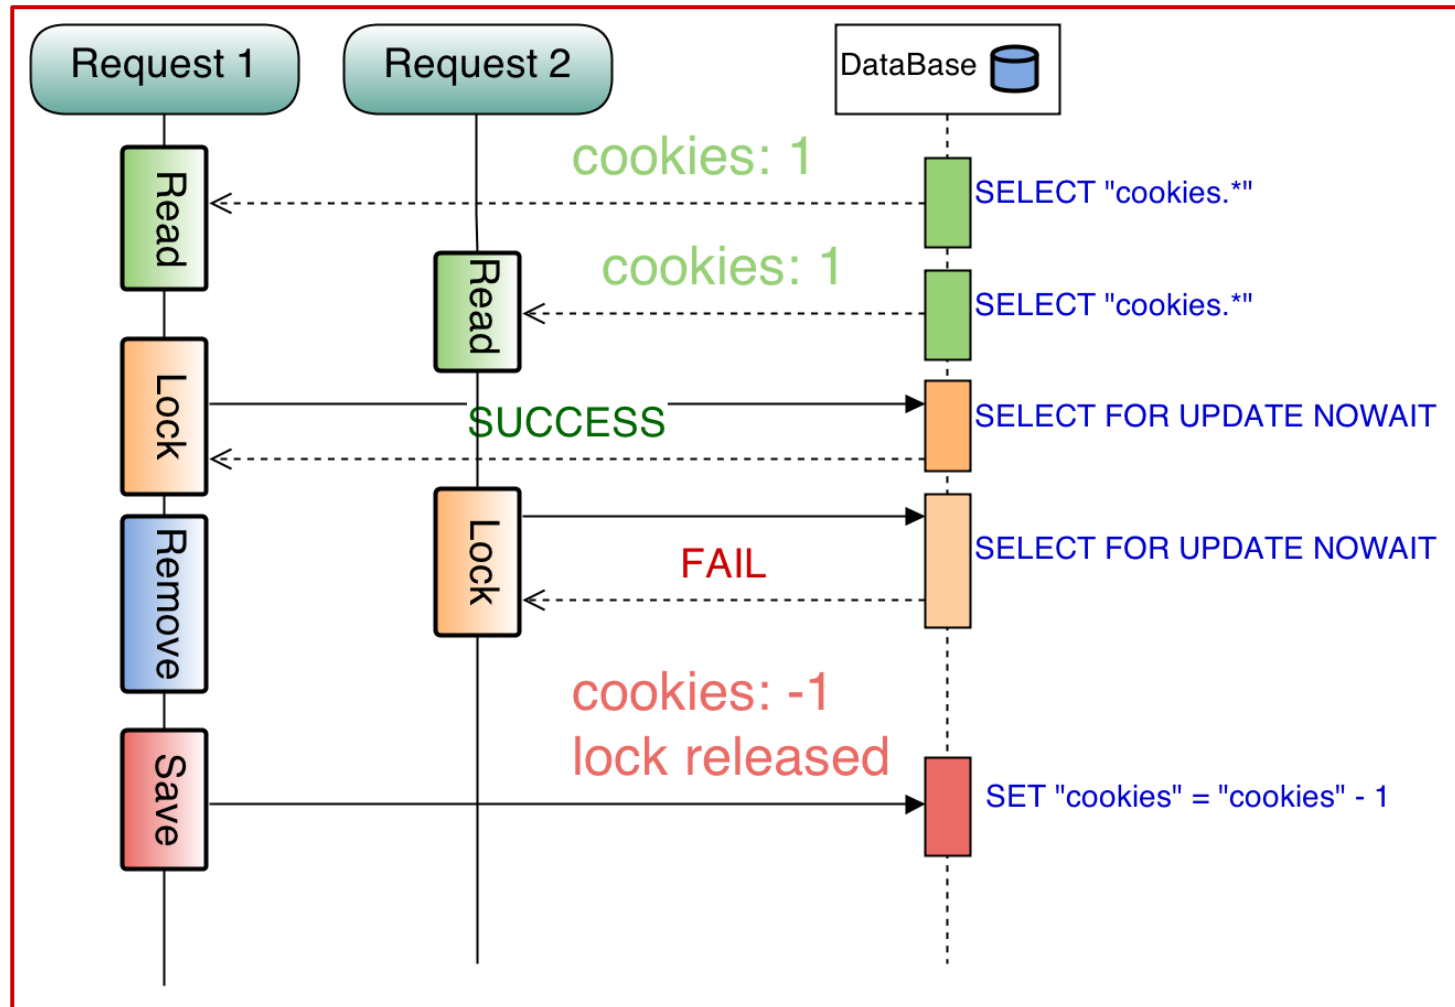
- easy to use
- all requests are processed

Cons:

- response time increases with concurrency requests
- deadlocks!

```
1  @stockpile = Stockpile.find_by_id(params[:id])
2  @stockpile.with_lock do
3    @stockpile.cookies += 1
4    @stockpile.save
5  end
```

Pessimistic locking (nowait)



Pessimistic locking (nowait)

Pros:

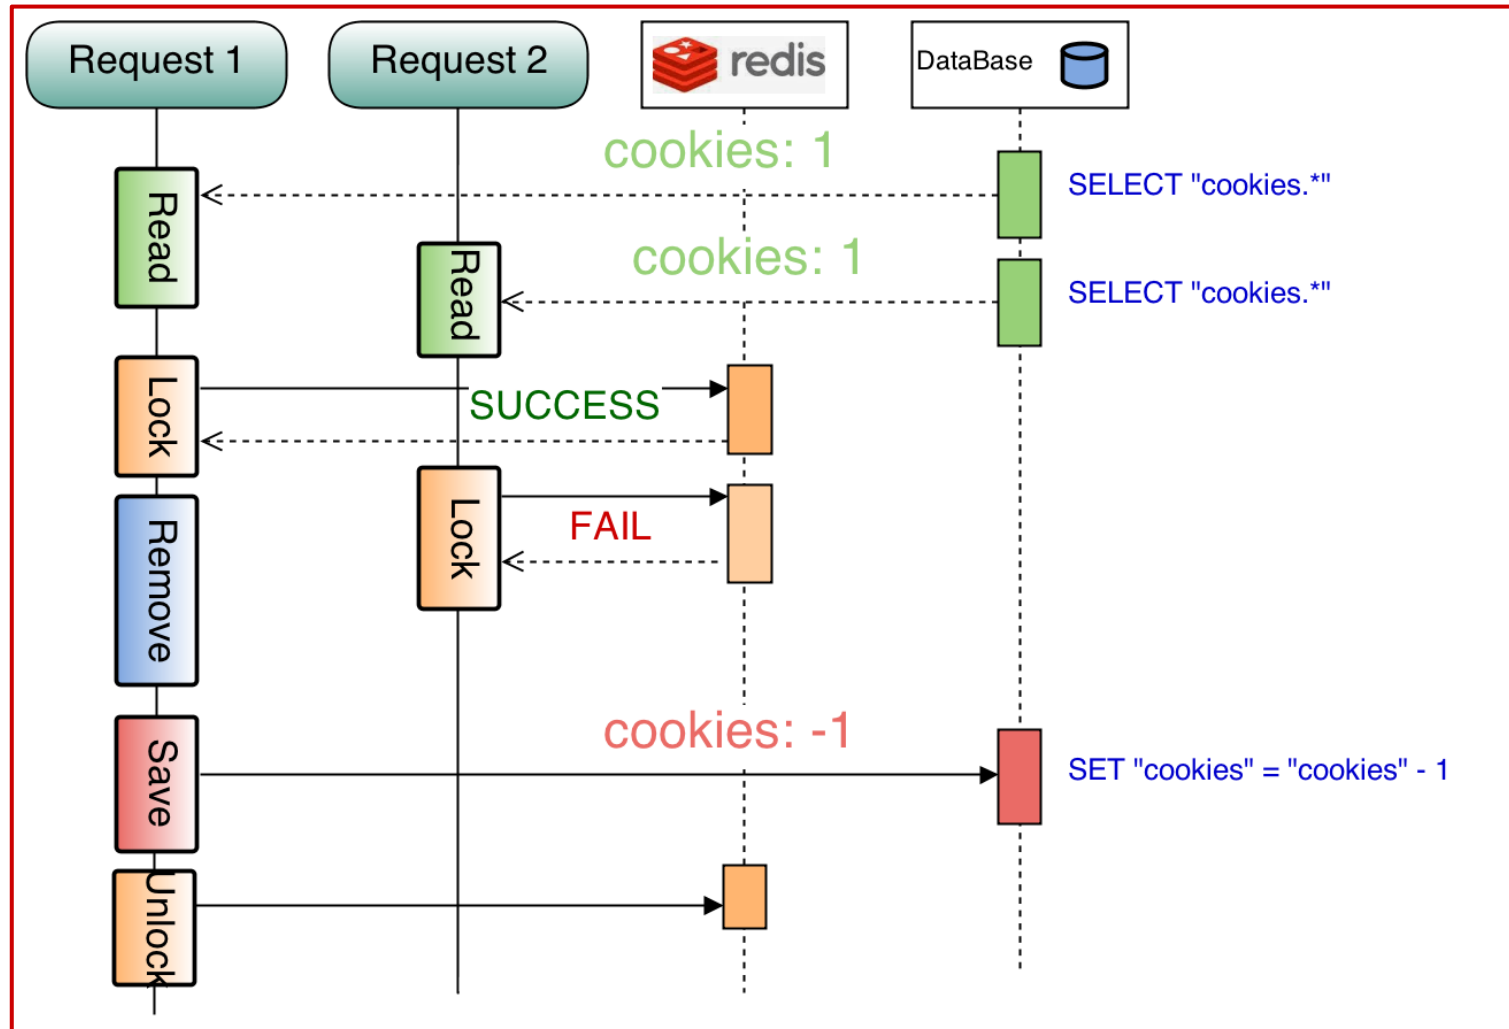
- no deadlocks

Cons:

- records are locked for read (unless share mode)
- need to rescue exceptions

```
1  #dont forget to handle exceptions
2  @stockpile = Stockpile.find(params[:id], lock: 'FOR UPDATE NOWAIT')
3  @stockpile.cookies += 1
4  @stockpile.save
```

Redis locking



Redis locking

Pros

- reduces database load
- more flexible (concentrate on synchronized code)

Cons

- easy to fail (data can be updated from other unlocked places)
- need to manually reload models

```
1  def lock(lock_id, &block)
2    return false unless REDIS.set(lock_id, true, {ex: 5.seconds, nx: true})
3    begin
4      yield
5    ensure
6      REDIS.del(lock_id)
7    end
8  end
```

Proper Redis Locking Pattern

Problem 3

Duplicated requests

“Мотороллер не мой!!!
Я просто разместил ОБЪЯВУ!”



Double Request Id

After page render:

```
1  after_filter :double_request_lock
2
3  def double_request_lock
4    redis.setnx(lock_key('double_request_id'), 'flag')
5  end
```

Before modification:

```
1  before_filter :check_double_request, only: :buy
2
3  def check_double_request
4    unless REDIS.del(lock_key('double_request_id')) == 1
5      redirect_to hunting_path
6    end
7  end
```

Double Request Id

Pros:

- easy to use
- invisible for user

Cons:

- not RESTful
 - doesn't provide 100% protection (in case of request queue)
-

Redis cooldown

Without magic:

```
1  before_filter :check_cooldown, only: :add
2
3  def check_cooldown
4    return if REDIS.set('add_cookies_cooldown', true, {px: 100, nx: true})
5
6    respond_to do |format|
7      format.json { render json: {stockpile: {cookies: @hunter.cookies}} }
8      format.html { redirect_to hunting_path }
9    end
10  end
```



COOL DOWN

Like a boss

Redis cooldown

Pros:

- easy to use
- invisible for user
- RESTful

Cons:

- doesn't provide 100% protection (in case of request queue)
-

We need to test this stuff!

It's difficult to simulate in development

It's impossible to test manually

Life is too short for manual testing



Model level testing

Helpers:

```
1  def run_in_thread(*args, &block)
2    Thread.new(*args) do |*args|
3      begin
4        yield(*args)
5      ensure
6        ActiveRecord::Base.connection.close
7      end
8    end
9  end
10
11 def several_threads(arguments = 4.times.map, &block)
12   running = arguments.map do |argument|
13     sleep(0.01)
14     run_in_thread(argument, &block)
15   end
16   running.each { |thread| thread.join }
17 end
```

Model level testing

Test example:

```
1  describe Concurrency::LockStrategies::Redis do
2    before do
3      @bucket = create(:bucket)
4      Cookable.stub(:change_testing_hook).and_return { sleep(0.1) }
5    end
6
7    it 'should be thread safe and execute only first call' do
8      several_threads([1, 2, 3]) do |amount|
9        bucket = Bucket.find_by_id(@bucket.id)
10       bucket.add(amount)
11     end
12
13     @bucket.reload.cookies.should == 1
14   end
15 end
```

Model level testing

Pros:

- easy to write
- possible to simulate any situation
- fast for all time execution


Cons:

- doesn't include full application stack
-

Integration tests with Capybara

Helpers:

```
1  def run_in_process(*args, &block)
2    fork { with_reconnect(*args, &block) }
3  end
4
5  def several_processes(arguments = 4.times.map, &block)
6    with_reconnect do
7      running = arguments.map do |args|
8        sleep(0.01)
9        run_in_process(args, &block)
10      end
11      running.each { |process_id| Process.waitpid(process_id) }
12    end
13  end
14
15  def with_reconnect(*args, &block)
16    begin
17      REDIS.client.reconnect
18      spec = Rails.application.config.database_configuration[Rails.env]
19      ActiveRecord::Base.establish_connection(spec)
20      yield(*args)
21    ensure
22      ActiveRecord::Base.connection.close
23    end
24  end
```



Integration tests with Capybara

Test example:

```
1  feature 'Concurrent getting from bucket' do
2    before do
3      @hunters = 3.times.map { create(:hunter) }
4      steal_bucket.update_attribute(:cookies, 10)
5    end
6
7    it 'should maintain total cookies amount' do
8      several_processes(10.times.map) do
9        login_hunter(@hunters.sample)
10       visit hunting_path
11       click_on 'get_steal_bucket_link'
12     end
13
14     total_cookies.should == 10
15   end
16 end
```


Integration tests with Capybara

Pros:

- involves whole application stack
- detects improper synchronization
- almost like in production

Cons:

- much slower
 - harder to find assertion criterias
-

Production level testing with Vagrant

Pros:


- simple to run development code in production.
- test database for simplified testing
- production database config for paranoiacs
- same integration tests as on local machine

Cons:

- slow
 - very slow
-


Testing gotchas

- write failing test before fixing
- results could be different on different configurations
- use edge conditions, avoid exact tests
- print statements as debug mode



Conclusion

The concurrent modifications **could be a problem**, but we can solve it, because we **can test it!**



Thank you!

Questions?
